


# Policy Manifold Search for Improving Diversity-based Neuroevolution

Nemanja **Rakicevic**  
Antoine **Cully**  
Petar **Kormushev**

Imperial College  
London



## Reinforcement Learning Problem Setting




Agent, defined by a policy ( $\pi$ ), interacts with the environment, by making observations ( $o$ ) and determining the best action ( $a$ ) to execute next.


Standard RL approaches learn a single optimal policy.

## Why Behaviour Diversity?

If variations in the environment are too great, or if situations occur, which have not been seen sufficiently during training - a single policy might not be able to solve the task anymore.



Robot damage recovery [1]



Encountering obstacles not seen previously [2]

Having **multiple policies**, exhibiting **diverse behaviours** in the environment, would be useful in these situations.


Main approaches in the literature:

- Skill-conditioned policies [3, 4]
- Quality-Diversity [1, 5]

## Behaviour descriptor

**Behaviour descriptor** quantifies a policy behaviour, by mapping the trajectory generated by a policy in an environment, to a **behaviour index (b)**, which defines the position of the corresponding policy parameters in a **policy collection**.

Policy parameters are added to the policy collection, if they exhibit diverse behaviours in the environment during evaluation.



## Policy Manifold Search


### Step 0) Initialisation

Start by randomly sampling a set of policy parameters  $\{\theta^*\}_k$

$$\{N(0, I)\}_k \rightarrow \{\theta^*\}_k$$

### Step 1) Evaluate and Add to the Policy Collection

The generated set of policies is evaluated in the environment.




### HYPOTHESIS

There exists a low-dimensional **manifold**, embedded in the high-dimensional **policy parameter space**, around which a **high-density of solutions** can be found.


### Step 2) Manifold Learning

Current policy collection is used to train the AE and obtain a latent representation of the policy parameter space.




### Step 3) Manifold Search

Sample the policy collection as in [3] and assess the reconstruction error of the selected sample, based on the current AE state.




Search either in latent or parameter space, depending on the value.

$$\mathcal{L}_{AE}(\theta) \geq \epsilon_{recn}$$




$$\mathcal{L}_{AE}(\theta) < \epsilon_{recn}$$




**Repeat Steps 1 - 3**, until stopping criteria are met.

## Experimental Evaluation


### Environments



Striker




Bipedal Walker





Bipedal Kicker

### Normalized observations



### Mixed-Scale observations



## Conclusions

- Search in the learned latent space (implicitly or explicitly) is better than in the original high-dimensional parameter space.
- Decoder Jacobian-based latent search scaling is crucial.
- Nonlinear representations are desired over linear representations.

## References

- [1] Cully, Antoine et al. (2015). "Robots that can adapt like animals." In: *Nature*
- [2] Ichnowski, Jeffrey et al. (2020). "Cloud-based motion plan computation for power-constrained robots." In: *Algorithmic Foundations of Robotics XII*
- [3] Eysenbach, Benjamin, et al. (2018). "Diversity is All You Need: Learning Skills without a Reward Function". In: *ICLR*
- [4] Hausman, Karol et al. (2018). "Learning an embedding space for transferable robot skills." In: *ICLR*
- [5] Cully, Antoine and Yiannis Demiris (2017). "Quality and diversity optimization: A unifying modular framework." In: *IEEE Transactions on Evolutionary Computation*